

1st October 2024

MSc in Software Engineering

Thesis Defense

# Improving Parallel and Concurrent Programming in FreeST

## STUDENT

Guilherme João Correia Lopes

## ADVISORS

Andreia Mordido

Vasco Vasconcelos



# Improve FreeST's practical parallel and concurrent capabilities!




## Analysis

Identifying difficulties and challenges concerning usual parallel and concurrent programming practices.




## Proposal

Elaboration of tools that address the identified challenges and improve the development experience.



## Integration

Development of tools into modules suited to the language characteristics.



### Data parallelism

Addressing embarrassingly parallel problems




### Futures


Asynchronous placeholders for pending results



- Data is divided into subsets and processed simultaneously by multiple processors.
- An embarrassingly parallel problem where minimal effort is required to separate the problem into multiple parallel tasks.


- The representation of a promise to deliver the value of an expression at some later time.





### Divide-and-conquer

Divide a problem, solve subproblems  
and merge solutions



### Streams

Efficient handling of continuous  
data processing

1. **Divide:** Break down a problem into smaller subproblems.
2. **Conquer:** Solve the subproblems recursively.
3. **Merge:** Combine the results into the solution.

- Flows of data elements that can be continuously and sequentially read from or written to.
- **Filters:** Functions that process data by maintaining endpoints to send and receive data elements within streams.
- **Pipelines:** Multiple filters linked in a sequence.

## MPI


- A message-passing specification.
- Facilitates data exchange between processes through cooperative operations.
- Harnesses data parallelism: optimal for embarrassingly parallel problems.

## ForkJoin

- A concurrent divide-and-conquer technique.
- Uses **fork** to divide and **join** to merge, which are analogous to futures.
- Implements a work-stealing scheduler.

## StreamIt

- A programming language purposely designed for efficient stream programming.
- Uses standard stream mechanisms such as **filters** and **pipelines**.
- Additionally, it provides **splitters**, which distribute data from one filter to other two.





## Session Types

- Focus on binary (two-party) sessions by well-defined protocols between endpoints.
- Guarantees that all agents involved in a communication strictly follow a protocol, ensuring reliable communication.
- **Compile-time!**



```
type ClientChannel = !Int;?Int  
type ServerChannel = ?Int;!Int
```



```
type ClientChannel = +{IsOdd: !Int;?Boolean,  
                      Succ : !Int;?Int}  
type ServerChannel = &{IsOdd: ?Int;!Boolean,  
                      Succ : ?Int;!Int}
```




- Focus on binary (two-party) sessions by well-defined protocols between endpoints.
- Guarantees that all agents involved in a communication strictly follow a protocol, ensuring reliable communication.
- **Compile-time!**



!?

```
type ClientChannel = !Int;?Int  
type ServerChannel = dualof ClientChannel
```


&+

```
type ClientChannel = +{IsOdd: !Int;?Boolean,  
                      Succ : !Int;?Int}  
type ServerChannel = dualof ClientChannel
```







## FreeST is not simple!



**Maturity:** FreeST is a new programming language in an early stage of development.




**Accessibility:** Its features and type system might impose difficulties for newcomers (e.g., linearity).




**Flexibility:** It has a limited set of tools, restricting its real-world utility in parallel and concurrent scenarios.

## How can we mitigate these challenges?




Three new modules, each providing user-friendly environments and abstractions for different parallel and concurrent programming concepts previously discussed.




## Parallel module

A parallel and concurrent programming environment that addresses data parallelism.




## Futures module

Implements futures and allows divide-and-conquer algorithms.




## Streams module

Provides a set of abstractions suited for stream programming.




## Parallel module

A parallel and concurrent programming environment that addresses data parallelism.



## Futures module


Implements futures and allows divide-and-conquer algorithms.



## Streams module

Provides a set of abstractions suited for stream programming.

## Parallel module: Parallel average example




```

type ParallelStream = ![Int];?Int;Wait


process : dualof ParallelStream -> ()
process c = let (xs, c) = receive c in
            c |> send (sum xs) |> close

parallelAverage : [Int] -> ParallelStream -> ParallelStream
                  1-> ParallelStream 1-> Int
parallelAverage xs w1 w2 w3 =
  let (xs, ys) = splitAt 3 xs in
  let (ys, zs) = splitAt 3 ys in
  let w1 = send xs w1 in
  let w2 = send ys w2 in
  let w3 = send zs w3 in
  let (x1, w1) = receive w1 in
  let (x2, w2) = receive w2 in
  let (x3, w3) = receive w3 in
  wait w1; wait w2; wait w3;
  (x1 + x2 + x3) / 9

main : Int
main = let w1 = forkWith process in
       let w2 = forkWith process in
       let w3 = forkWith process in
       parallelAverage[1,2,4,8,16,32,64,128,256] w1 w2 w3
  
```



**Repetitive**  
could benefit from  
some abstraction



**Unscalable**  
not ready to expand  
the workforce



**Complex**  
simple problem,  
unintuitive implementation

```

type ParallelStream = ![Int];?Int;Wait

process : dualof ParallelStream -> ()
process c = let (xs, c) = receive c in
            c |> send (sum xs) |> close




parallelAverage : [Int] -> ParallelStream -> ParallelStream
                  1-> ParallelStream 1-> Int
parallelAverage xs w1 w2 w3 =
  let (xs, ys) = splitAt 3 xs in
  let (ys, zs) = splitAt 3 ys in


  let w1 = send xs w1 in
  let w2 = send ys w2 in
  let w3 = send zs w3 in

  let (x1, w1) = receive w1 in
  let (x2, w2) = receive w2 in
  let (x3, w3) = receive w3 in
  wait w1; wait w2; wait w3;

  (x1 + x2 + x3) / 9

main : Int
main = let w1 = forkWith process in
       let w2 = forkWith process in
       let w3 = forkWith process in
       parallelAverage[1,2,4,8,16,32,64,128,256] w1 w2 w3
    
```

**Distribution****Broadcast****Scatter****Gathering****sum****Reduce****Gather**



## IV Parallel module: Implementation

Session types to outline the communication:

```
type ManagerStream = +{ Broadcast: ![Int]
, Scatter : ![Int]
, Gather : ?[Int]
, Reduce : ?Int
, Done : Wait} ; ManagerStream
```


```
type WorkerStream = dualof ManagerStream
```


List of endpoints to communicate with the workers:

```
data Comm = WNil () | Worker ManagerStream Comm
```

Initialize the manager-workers communication framework:

```
initialize : (Comm -> a) -> (WorkerStream -> ())
-> Int -> a
```






```

type ParallelStream = ![Int];?Int;Wait

process : dualof ParallelStream -> ()
process c = let (xs, c) = receive c in
            c |> send (sum xs) |> close

parallelAverage : [Int] -> ParallelStream -> ParallelStream
parallelAverage xs w1 w2 w3 =
  let (xs, ys) = splitAt 3 xs in
  let (ys, zs) = splitAt 3 ys in
  let w1 = send xs w1 in
  let w2 = send ys w2 in
  let w3 = send zs w3 in
  let (x1, w1) = receive w1 in
  let (x2, w2) = receive w2 in
  let (x3, w3) = receive w3 in
  wait w1; wait w2; wait w3;
  (x1 + x2 + x3) / 9

main : Int
main = let w1 = forkWith process in
       let w2 = forkWith process in
       let w3 = forkWith process in
       parallelAverage[1,2,4,8,16,32,64,128,256] w1 w2 w3
  
```



```


import Parallel

worker : WorkerStream -> ()
worker c = let (xs, c) = wscatter c in
           c |> wreduce (sum xs) |> wdone

manager : [Int] -> Comm -> Int
manager xs comm = let xsl = length xs in
                  let comm = mscatter xs comm in
                  let (xs, comm) = mreduce (+) 0 comm in
                  mdone comm; xs / xsl

main : Int
main = let xs = [1, 2, 4, 8, 16, 32, 64, 128, 256] in
       initialize (manager xs) worker 3
  
```

Note: The `manager` selects operations via `m` prefix; the `worker` matches with `w` prefix.




```

.
.
.

nfoldl : Int -> (Int -> Int -> Int) -> Int -> [Int] -> Int
nfilter : Int -> (Int -> Bool) -> [Int] -> [Int]
nzipWith : Int -> (Int -> Int -> Int) -> [Int] -> [Int] -> [Int]


import Parallel

main : Int
main = let xs = [1, 2, 4, 8, 16, 32, 64, 128, 256] in
      nfoldl 3 (+) 0 xs / length xs
  
```




## Parallel module

A parallel and concurrent programming environment that addresses data parallelism.



## Futures module


Implements futures and allows divide-and-conquer algorithms.



## Streams module

Provides a set of abstractions suited for stream programming.

## How can we retrieve the result of an asynchronous computation?



Could we avoid this hassle? **YES!**

### Futures

A mechanism that wraps the `fork` primitive by abstracting the creation of an extra channel, simplifying retrieving the result of an asynchronous computation.

Launch an asynchronous computation:



```
future : (( ) -> a) -> ?a;Wait
```

Retrieve the result of the computation:

```
block : ?a;Wait -> a
```

Asynchronously delay a computation:

```
delay : (( ) -> a) -> (( ) -> b)  
-> ?b;Wait
```




Launch an asynchronous computation:

```
future : (( ) -> a) -> ?a;Wait
```

Retrieve the result of the computation:

```
block : ?a;Wait -> a
```



**Example**


The Fibonacci sequence defines each number as the sum of the two preceding ones.

$$F_0 = 0, \quad F_1 = 1,$$

$$F_n = F_{n-1} + F_{n-2}$$


**Implementation with futures**

```
pFib : Int -> Int
pFib n | n == 0    = 0
       | n == 1    = 1
       | otherwise = let f1 = future (\_:() -> pFib (n - 1)) in
                     let f2 = future (\_:() -> pFib (n - 2)) in
                     block f1 + block f2
```




## Parallel module

A parallel and concurrent programming environment that addresses data parallelism.



## Futures module

Implements futures and allows divide-and-conquer algorithms.



## Streams module

Provides a set of abstractions suited for stream programming.


Session types and channels allow for a natural way of writing and handling streams in FreeST!

#### Example of a stream of integers

```
type IntStream = +{More: !Int; IntStream, Done: Wait}
```

#### Example of a stream program

```
main : ()  
main = let (w, r) = new @IntStream () in  
      fork (\_:(() -> intSource w);  
            intPrinter r)
```



Session types for streams:


```
type OStream = +{More: !Int;OStream, Done: Wait}  
type IStream = dualof OStream
```

Basic operations on streams:

```
sends : Int -> OStream -> OStream  
waits : OStream -> ()  
forward : IStream -> OStream 1-> OStream
```

List related operations on streams:

```
fromList : [Int] -> OStream -> ()  
toList : IStream -> [Int]
```



**Definition**


**Splitters** distribute data from a stream between two streams.


Divide a stream into two:

```
splitSDup : IStream -> OStream 1-> OStream 1-> ()
splitSAlt : IStream -> OStream 1-> OStream 1-> ()
splitSWith : (Int -> Bool) -> IStream -> OStream
           1-> OStream 1-> ()
```

Merge two streams into one:

```
joiner : IStream -> IStream 1-> OStream 1-> ()
```







`fromList`  
`splitSWith`  
`foldlS`

`mapS : (Int -> Int) -> IStream -> OStream 1-> ()`  
`filterS : (Int -> Bool) -> IStream -> OStream 1-> ()`  
`foldlS : (Int -> Int -> Int) -> Int -> IStream -> Int`

**import Streams**  
`foo : [Int] -> (Int, Int)`  
`foo xs = let i1 = forkWith (fromList xs) in`  
 `let (i2, i3) = forkWith2 (splitSWith even i1) in`  
 `(foldlS (+) 0 i2, foldlS (+) 0 i3)`

`main : (Int, Int)`  
`main = foo [1, 6, 23, 82, 34, 2, 5, 44, 302, 48, 17, 3]`




**Target**

We only submit the **Parallel** module to this analysis.

**Goal**

LoC comparison between examples implemented with FreeST and with the **Parallel** module.



**Thank you!**

**Questions**

**Answers**

**R-Q&A**